



NLP IN PYTHON

WWW.HACKYALE.COM



NLP IN PYTHON


WEEK 1
CLASSIFYING TEXT

SOURCES

- Natural Language Processing with Python, by Steven Bird, Ewan Klein, and Edward Loper. O'Reilly Media, 978-0-596-51649-9.
- MOOC: Natural Language Processing with Dan Jurafsky and Christopher Manning (Stanford, Coursera)

REVIEW

NLTK BOOK: CORPUS CHART



CREDIT: NLTK BOOK

INAUGURAL CORPUS

```
from nltk.corpus import inaugural  
import matplotlib
```

```
from nltk import Text  
inaug = Text(inaugural.words())  
inaug.collocations()  
inaug.concordance('freedom')  
inaug.dispersion_plot(['freedom', 'government',  
                        'liberty', 'hope'])
```

WHAT IS TEXT CLASSIFICATION?

TEXT CLASSIFICATION

- Building a model on some features of a language
- Assigning categories to unseen documents, given the ones you've seen
- Documents can mean anything:
 - Emails, news articles, tweets, etc.
 - Spam detection, article relevance, sentiment analysis, etc.

TEXT CLASSIFICATION

- On different levels
 - By word, phoneme, author, genre
- **Identifying patterns and making predictions**

TRAINING

- Supervised
 - Hand-labeled documents (i.e. movie reviews with pos. or neg.)
- Unsupervised
 - Completely unlabeled
- Semi-supervised
 - Mixture of both

TRAINING

- **Supervised**

- Hand-labeled documents (i.e. movie reviews with pos. or neg.)

- **Unsupervised**

- Completely unlabeled

- **Semi-supervised**

- Mixture of both

BAYES RULE

$$P(A|B) = \frac{P(B|A)P(A)}{P(B)}$$

NOISY CHANNEL MODEL

- There is some process that introduces noise to an underlying “true” form
- Misspelling:
 - The “true” spelling of a word goes through some noisy process that produces a misspelled word
 - Edit distance, position in word, proximity of keys on keyboard, homophones, confusion of left and right fingers

BAYES RULE

$$P(cat|obs) = \frac{P(obs|cat)P(cat)}{P(obs)}$$

cat = category

obs = observation

NLP EXAMPLES

Speech Analysis:

$$P(/keep/|"keep") = \frac{P("keep" | /keep/) P(/keep/)}{P("keep")}$$

Spelling Corrector:

$$P(the|thew) = \frac{P(thew|the)P(the)}{P(thew)}$$

DEFINITIONS

$$P(/keep/|"keep") = \frac{P("keep" | /keep/)P(/keep/)}{P("keep")}$$

$P("keep" | /keep/)$ - Noisy process

$P(/keep/)$ - Prior probability of the lexicon

NAIVE BAYES INFERENCE

- Independence assumptions
 - Noisy model can be interpreted without context
- Our case:
 - Labeled categories
 - Look at inputs separately
 - Independence assumption applies to words in the topics

LANGUAGE MODELS

LANGUAGE MODEL

- N-gram
 - Simplest and “dumbest” model
 - Just counts words and their physical relation to each other in the document
 - Or phonemes, or letters, etc.
- Bag of words
 - Google n-grams example

N-GRAMS

- Applications
 - Author identification
 - Federalist Papers
 - Genre identification

N-GRAMS

Unigram (simple count):

“This movie is great!”

{‘this’: 1, ‘movie’: 1, ...}

Bigram:

“# This movie is great!”

[(‘#’, ‘this’), (‘this’, ‘movie’), (‘movie’, ‘is’), ...]

N-GRAMS

Trigram:

“# # This movie is great”

[(' ', ' ', 'this'), (' ', 'this', 'movie'), ...]

4-grams, 5-grams, etc, etc.

Trigrams are usually good for English

N-GRAMS

- You have to strike a balance between
 - Capturing context (large n)
 - Creating a generalizable model (small n)
- Flaws:
 - Doesn't model long-distance dependencies
 - “The movie that I saw with my friends on Wednesday was great.”

NLP EXAMPLE (BIGRAMS)

$$\begin{aligned} &P(\text{"I hate cats"}) \\ &= P(I \mid \text{START}) P(\text{hate} \mid I) P(\text{cats} \mid \text{hate}) \end{aligned}$$

How do we measure this?

Relative Frequency Estimation (i.e., counting)

$$P(\text{hate} \mid I) = \text{count "I hate"} / \text{count "I"}$$

PROBABILITY OF A SENTENCE

INDEPENDENCE ASSUMPTION

Joint prob of a,b = prob of a given b and the prob of b

$$P(a, b) = P(a \mid b) P(b)$$

Markov Chain:

$$P(a,b,c,d) = P(a \mid b,c,d) P(b \mid c,d) P(c \mid d) P(d)$$

Independence Assumption (for Bigrams):

$$P(a,b,c,d) = P(a \mid b) P(b \mid c) P(c \mid d) P(d)$$

N-GRAMS

```
# -*- coding: utf-8 -*-  
from nltk import bigrams, ConditionalFreqDist  
from nltk.corpus import gutenber  
  
bi = bigrams(gutenberg.words(fileids="austen-  
emma.txt"))  
cfd = ConditionalFreqDist(bi)
```

PROBLEMS WITH N-GRAMS

- New words will have 0 counts
 - If your corpus only recognizes this sentence:
 - “This is great.”
 - And you get:
 - “This is amazing.”
 - $P(\text{this} \mid \text{START}) P(\text{is} \mid \text{this}) P(\text{amazing} \mid \text{is})$
 - The probability of this sentence is 0.

SMOOTHING

- Re-distributes probability mass from recognizable words to unrecognized words
- Ideally, pulls more mass from the n-grams we're not certain about (i.e., those with low counts)
- Basic: Laplace smoothing
 - +1 to every 0 count then normalize
 - Add-Lambda: use something other than 1!

SMOOTHING

- More advanced: Good-Turing
 - Give the probability of the 1-counts to the 0-counts, the 2-counts to the 1-counts, the 3-counts to the 2-counts, etc., etc.
 - Then normalize
 - You have to stop eventually
 - Sometimes you'll have categories that don't exist (i.e., there are no words counted 18 times).
 - What do you do for the 17-count words?

BASIC TEXT CLASSIFICATION

IMPORTANT STEPS

- Define a features function
 - {'feature' : count/bool, 'feature2' : count/bool}
- Label words in your corpus with their categories
 - [(category, word), (category, word), (category, word)]
- Create a features set
 - [(list_of_features, category), ...]

IMPORTANT STEPS

- Create a classifier object
- Remember that features are just a dict of the important aspects of some word/document
- Jack => {'last_letter' : k}

NAMES CORPUS

- 8000 names, separated into male and female
- Think of appropriate features
- Overfitting is a problem
 - Especially with a small data set

MOVIE REVIEWS

- Mess around with `corpus.movie_reviews`
- See `sentiment.py`
- Play around with numbers in NLTK Book example
 - (i.e., size of test set vs. training set, # of most frequent words we consider relevant)

ACCURACY VS. F-SCORE

	<u>True</u>	<u>False</u>
<u>Chosen</u>	True Positive	False Positive
<u>Not Chosen</u>	False Negative	True Negative

Accuracy:

How often do we choose correctly?

(i.e. True / False)

True Negative might be HUGE compared to True Positive, skewing the results.

ACCURACY VS. PRECISION & RECALL

	<u>True</u>	<u>False</u>
<u>Chosen</u>	True Positive	False Positive
<u>Not Chosen</u>	False Negative	True Negative

Precision:

Of the things our model selects, how many belong in the category?

True Positive / False Positive

ACCURACY VS. F-SCORE

	<u>True</u>	<u>False</u>
<u>Chosen</u>	True Positive	False Positive
<u>Not Chosen</u>	False Negative	True Negative

Recall:

Of the things that belong to the category, how many does our model select?

True Positive / False Negative

BALANCED F1-SCORE

Combines Precision and Recall values:

$$\frac{2PR}{P + R}$$

ACCURACY

	True	False
Selected	True Positive	False Positive
Not Selected	False Negative	True Negative

Precision: Of the things the model selects, how many truly belong in that category.

Recall: Of the things in that category, how